

Rodrigo Neves

STAFF ENGINEER

✉ rodrigo.f.l.neves@gmail.com | 🏠 rnev.es | 📷 RodrigoNeves95 | 🌐 rodrigoflneves

“Only a few know how much one must know to know how little one knows”

Summary

Product Engineer, motivated by taking ideas from discovery to production, building is what drives me. I thrive in high-autonomy environments with a bias to action, owning ambiguous problems end-to-end and optimizing for fast, high-quality iteration (fail fast, learn fast). Generalist by heart, love to learn about new stuff, which is a curse and a blessing. AI adopter, and a firm believer that, in today's world, any developer not using the tools out there is going to be left behind. The caveat is, fundamentals matter more than ever, and developers critical thinking and deep understanding about building software matters more than ever.

Work Experience

Syrup Tech (acquired by Anaplan, 2025)

New York, United States

STAFF ENGINEER

Jan 2024 - Present

- Led the transition from a single-tenant platform to a multi-tenant architecture, shaping the migration plan and aligning product and engineering on scope, sequencing, and tradeoffs.
- Wore the Solution Architect hat for early migrations, partnering directly with founders and customers to translate requirements into technical design docs, rollout plans, and successful cutovers.
- Migrated and onboarded 15+ tenants while supporting customers on both platforms, leveraging deep end-to-end product knowledge to unblock teams and keep the legacy platform stable.
- Drove reliability, performance, and cost improvements across the stack by identifying bottlenecks and optimizing data flows and query patterns in Snowflake/dbt and Postgres.
- Co-led a greenfield “Buying” product POC from discovery to production in 3 months, running 20+ customer/expert interviews, shipping 2 beta pilots, and owning end-to-end data/backend integration of forecasting outputs into an interactive recommendation workflow.

Syrup Tech

New York, United States

SENIOR FULL-STACK ENGINEER

April 2022 - Jan 2024

- Founding engineer #2; led development of the core optimization/intelligence layer that generated business-critical recommendations, reducing runtime from 8+ hours to under 1 hour and automating previously manual workflows.
- Built product features end-to-end across data/optimization pipelines, backend APIs, and frontend UX, bringing a product engineer mindset with high ownership and customer empathy.
- Owned production operations for critical services during rapid scale, improving debuggability and speed-to-recovery through quality-of-life tooling for re-runs and faster iteration.
- Helped scale the company from pre-seed to Series A (Accel-led, 16M+), growing from 3 customers to 15+ and supporting the shift from customer-first delivery to building a platform for the next 100+ customers.

Daredata Engineering

Lisbon, Portugal

MACHINE LEARNING ENGINEER

June 2021 - Present

- Migrated cross-selling models from R into Python, further improving them with new feature engineering and models architectures and implementing distributed hyperparameter optimization on Spark with Hyperopt. Pushed best practices from MLOps perspective to easily operationalize the newly developed models, using Mlflow, Databricks, and AzureML on Azure platform.
- Built and deployed an MLOps based pipeline to deploy and schedule GPU-ready models on Google Cloud Platform using a Kubernetes cluster and Airflow. The main use-case was to run several deep learning NLP models to populate an Elasticsearch database with documentation summarization and entity recognition/extraction.
- Planned and developed a proof of concept solution for a big utility company to create an automatic and interactive solution to calculate solar potential through satellite images. Implemented a computer vision pipeline using Pytorch and OpenCV to generate all the artifacts and data needed to integrate with the GIS-based systems. As a second phase, a machine learning model was developed to calculate the probability of a customer buying solar panels, based on the output of the first phase.

Jungle AI

Lisbon, Portugal

SOFTWARE/MACHINE LEARNING ENGINEER

May 2018 - June 2021

- Researched and implemented Jungle's core approach to predictive modelling with an emphasis on recurrent models and time series data implementing from scratch most of the modules using Pytorch.
- Active member of 5+ projects within the energy and heavy industry sector with the main purpose of showing Jungle's predictive solutions. This work would evolve to gather, process and analyze data to apply Jungle's predictive maintenance pipeline and presenting the result to the clients.

- Designed, developed and build internal and production-ready software and deployments. Tasks like ORM definition for production databases, software to process essential high-volume data analysis or the development of a backend infrastructure to manage production deployments using FastAPI, among others.
- Designed and implemented a database architecture (Postgres and TimescaleDB) refactor to allow for a scalable, high-performing and cost-effective solution.
- Built an automated data collection and cleaning workflow (AWS S3 and Prefect) to process daily 2k+ sensors with a millisecond frequency.
- Built a CRUD RESTful API in Node.js for TimescaleDB and PostgreSQL database to be used by Jungle product.
- Implemented, deployed and maintained several high-performing, scalable and distributed workflows using Argo and Kubernetes making use of several CI/CD pipelines to help on automation.
- Designed, built and managed internal dashboards and data warehouse solutions using InfluxDB and Grafana to allow for high-volume data analysis.

Jungle AI

Lisbon, Portugal

- MASTER THESIS INTERNSHIP
- August 2017 - May 2018
- Had the opportunity to do my master thesis work in a professional environment. The main research goal was to study, implement and test the performance of different deep learning models and architectures applied to time-series problems, specifically in the area of predictive maintenance and forecasting. This allowed me to be an active member of Jungle’s daily projects to establish their technology.

Technical Skills

Programming Language	Python, TypeScript, SQL
Machine Learning Stack	PyTorch, Scikit-learn, XGBoost, LightGBM, MLflow
Data Analysis Stack	Pandas, NumPy, PySpark, Dask, OpenCV, Polars
Workflow Orchestration	Docker, Kubernetes, Argo, Prefect (Cloud), Airflow
Data	Snowflake, dbt, Databricks PostgreSQL, TimescaleDB, InfluxDB
, Cloud	GCP, Azure, AWS cloud (ECS, S3, EventBridge, SFTP, SQS, SNS, Lambda, etc.), Kafka, Redis
Software development	Git, CI/CD, Jira, Distributed Computing, IaC (Pulumi, Terraform, etc.)

Education

Insituto Superior Tecnico - University of Lisbon

Lisbon, Portugal

- MSC. IN ELECTRICAL AND COMPUTER ENGINEERING
- Sept. 2012 - June. 2018
- Masther thesis - "An overview of deep learning strategies for time-series prediction". Work done in collaboration with Jungle AI.